Nightwise — Full App Review / Debriefing

Expanded reconstruction of the complete review so far. This version preserves the detail, caveats, examples and implementation concerns rather than reducing the review to short bullets.

This document consolidates the review exactly as developed through the testing conversation. The intention is not simply to list bugs: where an observation could have more than one explanation, that uncertainty is preserved. The important distinction is between what was directly observed, what is suspected, and what still needs verification. At the same time, the non-negotiable product requirements remain clear: puzzles must be correctly configured, progress must be truthful and persistent, practice must not repeat incorrectly, and coaching must be context-aware.

1. “Here’s Your Briefing”

The “Here’s Your Briefing” option did not feel relevant or necessary in the current flow. Recheck whether it genuinely adds value or whether it is simply another layer competing with the actual coaching experience. The voice used in “Here’s Your Briefing” was also something to flag and re-evaluate.

2. Voice selection

It was not clear whether the voice selected in the voice picker was actually being applied. The selected voice should be re-verified by testing the voice that subsequently speaks in the relevant sections. The picker, saved state and actual playback should all agree.

3. Chess move interaction

Clicking a pawn correctly displayed its available moves. However, clicking the same piece again did not hide those available moves. Clicking another piece also left the previous pawn’s available-move indicators visible. Selection state therefore needs to be handled cleanly: selecting another piece should update the board to the new piece, and clicking the selected piece again should behave predictably rather than leaving stale indicators behind.

4. Today’s Five Minutes

The opponent’s move was not shown where it appeared to be needed. Verify opponent moves consistently in puzzles and similar sections so the user can understand the complete situation rather than only their own side of the position.

Keep the small Mindset tab on the right, above “I’m Ready”. Remove the text above it — “It’s the middle game, we have two pawns down” — because that information should not be sitting separately above the Mindset control.

Under “Puzzle: Trapping a Piece”, the three boxes read “Rated 1428”, “You are White”, and “Middle Game”. The “Rated 1428” item was unclear: it was not obvious why it was there or whether that rating was actually correct. Verify its source and whether it is useful to the user.

5. After “I’m Ready”

Remove the instruction “Click a piece, then the square to move it to.” The user wanted the actual coaching to remain: “One of their pieces has no safe square at all. Attack it.” The space after “I’m Ready” should be used for meaningful tactical guidance rather than repeating interface instructions.

6. Puzzle hints

After roughly 3–4 consecutive mistakes — with the exact threshold still to be determined — the coach should offer a small hint. The hint should guide the user towards the relevant piece, direction or idea without directly revealing the answer.

Repeated mistakes should make the coaching more helpful. Avoid repeatedly playing the same error sound and repeating exactly the same coach line. The system should recognise that the user is stuck and provide progressive guidance.

7. Puzzle completion and progress feedback

Progress information such as “0 of 2 right first time”, “0 positions learned for good”, and later “1 of 3, 0 right first time” was not sufficiently clear in context.

There was no obvious celebration or completion state, so the user did not realise that all three puzzles had been completed. After clicking Next, the experience returned to the first puzzle without clearly communicating that the session had ended.

The final puzzle should produce an unmistakable completion state and a clear transition. The user should know that the session is complete, what they achieved, and what happens next.

8. Guidance for new concepts and tactics

It was unclear whether the opening puzzle was introducing a new concept. Whenever a new tactic, concept or learning method is introduced — especially if it differs from ordinary gameplay — explain what is being taught and why it matters.

The coach should guide the user on which pieces to move, where and when appropriate, and how the idea would actually be used in a real game. The user should not have to infer the entire learning objective from the board alone.

9. Puzzle variety and repetition

Check whether “Your Game” puzzles are repeating every time. The app should select different suitable puzzles from the user’s games, whether from the same game or different games, rather than repeatedly serving the exact same position as fresh practice.

Variety should not come at the expense of relevance: the content should still be curated around the intended learning objective.

10. Practice-session layout

Remove the explanatory text above the Mindset tab as described earlier.

Under “Your Game vs Geetha Sharvesh”, there were too many tags: date, Rapid, You Were White, You Lost, Move 13, Middle Game. Keep only the most useful 3–4 pieces of metadata. Move 13 and Middle Game were not useful enough to justify the space.

Also remove the explanatory text below the practice section: “Your own positions”, “The move before a big mistake”, and “Puzzles on the same habit”. The user did not understand these labels and they took up space without making the practice objective clearer.

11. Game-position labels

Labels such as “Your Game vs M. Onis Irfer — Move 7, Still the Opening” and “Your Game vs Geetha Sharvesh — Move 13, The Middle Game” were unnecessarily repetitive. Simplify the presentation and avoid repeating the game title when the user is already inside the relevant game/context.

12. Coach briefing placement

The text displayed above the Mindset tab was what the coach actually spoke. Important briefing content should therefore be moved into the coach pop-up rather than appearing as a separate block on the page.

For example, instead of the generic introduction “Here’s the situation. Read it, then press ‘I’m Ready’”, the pop-up should contain the actual useful briefing. Not every piece of information has to be inside the pop-up, but the core coaching message should be delivered in the place where the user expects the coach to speak.

13. Arjun’s speaking speed

Arjun should speak more slowly if the recommended voice settings are applied. His current delivery can be too fast for instructional content, where comprehension matters more than conversational speed.

14. Progress saving

Progress in Today’s Five Minutes, Practice and other sections must persist. Switching tabs currently appeared to wipe or reset progress. Navigating between tabs must not erase completed work, current progress or relevant session state.

15. Daily refresh and curated practice

Practice content should refresh daily. If a question has already been attempted, it should not simply return as a fresh question. Only appropriate unattempted content should carry over when necessary.

Practice should remain curated. Check whether there is currently a deficit of suitable practice questions and ensure there is enough appropriate material without falling back to repeated puzzles. Daily refresh must coexist with saved progress.

16. Progress tab and badges

The Progress tab should primarily reflect actual Chess.com progress, while also tracking progress made inside Nightwise. Most badges may come from Chess.com, with some coming from Nightwise, but the two sources should not be presented in a misleadingly merged way.

17. Games section and key moments

When opening a game in the Games section, too many key moments and variations were visible simultaneously on the right. The content itself is useful, but showing too much text at once is overwhelming and could make a user close the tab.

The goal is to minimise what is immediately visible without reducing what the system can offer. Organise the information into specific, clear, engaging pieces so the user can focus on one thing at a time.

18. Game-review onboarding and Maya

When opening Game Review, Maya did not welcome the user immediately. She only appeared after the user attempted a key moment. Maya should welcome the user as soon as Game Review opens.

She should provide a short, voice-based briefing and recommendation, then guide the user through one key moment at a time. The total number of key moments can be shown, but only one should be presented for active work at a time.

The initial briefing should primarily be delivered through voice rather than a large block of text. The user indicated that more voice feedback could follow as testing continued.

19. Maya’s speaking speed

Maya should also speak more slowly. Her current speed is too fast for comfortable instructional listening.

20. Confetti and congratulations

The confetti and congratulations after correctly completing “Try It Again” worked as expected. This behaviour should be retained rather than removed while other completion feedback is improved.

21. Opening Tree usability

The Opening Tree was confusing for a beginner. A professional player might understand the notation, but the user struggled to interpret entries such as “1. e4” and “1... e5”, as well as the dropdowns and the overall tree at a glance.

The user was willing to spend time learning it because they understand the potential of the app, but a new user may simply skip the feature and never return. The Opening Tree therefore needs to become intuitive and beginner-friendly.

Explain what the notation means, how to navigate the tree, how to interpret branches and what the feature is actually helping the user learn. The goal is not to remove the professional functionality, but to make the first-use experience understandable.

22. Zara’s pointing indicator

It was unclear whether Zara’s finger/pointer was pointing to the correct plan or option. When clicked, it appeared to point somewhere else, and when minimised it appeared somewhere different again.

Verify the pointer against the actual UI target. If the option is already visible and the pointer does not add meaningful information, it may be better to remove it rather than risk misleading the user.

23. Welcome and onboarding across tabs

Every major tab or section needs a proper beginner-friendly introduction. Users should not be expected to immediately understand unfamiliar features. The exact implementation can be refined later, but the principle is that the app should guide the user through unfamiliar functionality rather than making them work it out alone.

24. “Winning with the Queen” puzzle

The user initially thought the objective was checkmate rather than stalemate. The puzzle ended after roughly 2–3 moves without checkmate; the user’s queen was captured and the game ended, yet the app marked the result as correct and displayed Next. There was also no confetti.

Importantly, the user recognised that this observation might not identify the true underlying problem. The puzzle could have had a different intended objective, or the objective/label might have been wrong. Therefore, do not automatically conclude that the puzzle itself is broken. Investigate the puzzle’s intended objective, configuration, accepted line and label.

Regardless of which cause is found, the final requirement is firm: no incomplete or incorrectly configured puzzles should reach users.

25. Chessboard capture indicators

When a piece can capture an opponent’s piece, show a small grey circle on the opponent’s piece. The current board indicates ordinary available movement squares but does not clearly distinguish captures. Capture availability should be visually clear.

26. “Winning with the Queen” inconsistent progress

The user saw “3 of 4” in one place, “4 right first time” elsewhere, and the Victory Coach said “All done, 4 of 4 right first time. That’s a new weapon for your games” while one puzzle still remained.

The user suspected that revisiting or completing a puzzle twice may have caused it to count twice. Verify unique puzzle identity, attempts, first-attempt correctness, completion state and every displayed total. All progress indicators, dots, completion messages and totals must agree.

27. General puzzle validation and review accuracy

A user's observation can reveal a symptom without necessarily identifying the underlying bug. A puzzle that feels incomplete could have a different intended objective, an incorrect label, an incorrect solution or incomplete configuration. Investigate the possibilities before deciding what the problem is.

However, the non-negotiable requirement is that the final product contain no incomplete or incorrectly configured puzzles, no misleading labels, no repeated practice content and no broken refresh/progress behaviour.

28. Puzzle refresh and progress integrity

Attempted puzzles must not reappear as fresh attempts. Progress must be saved. Completing or revisiting the same puzzle must not inflate progress. The system needs to track unique puzzles and actual attempts rather than simply counting successful clicks or completions.

This integrity rule applies across the entire app, not only to the specific puzzle where the issue was observed.

29. Incorrect “Right First Time” counting

After four puzzles, the last page showed “4 of 4” while a separate indicator showed “5 right first time”. Repeating puzzles then raised the count to six. This strongly suggests that successful attempts were being counted instead of unique puzzles solved correctly on the first attempt.

“Right First Time” must mean exactly that: each unique puzzle can contribute at most once, and only if the user solved it correctly on their first attempt. Repeating it later must not increase the number.

The same underlying truth must drive the progress dots, completion messages, totals and coach statements.

30. Badges tab clarity and achievement criteria

The badge presentation was visually attractive, with golden/brown/other coloured rings, but there was no legend explaining the colours. Add a small colour-coded legend, ideally in the top-right.

The user could not easily tell which badges were actually completed. The badge progress also appeared too easy — for example, 17 of 21 — potentially because many existing Chess.com games already satisfy the requirements.

A badge should count as fully achieved only when its actual requirements are met or its maximum level has been reached. Clearly distinguish progress towards a badge, the current level and full completion.

The current presentation did not strongly motivate badge pursuit because the user could not easily tell what remained to be done.

31. Voice settings status and API key

The Settings → Voice section said “Listening is on” even though the user had not configured an API key. That status is misleading.

The UI must accurately reflect whether an API key is configured and whether listening is actually available. Do not display a positive feature status merely because the feature exists in the interface.

32. API-key setup screenshots

Add screenshots showing the API-key setup process. They should make configuration straightforward for a user who has never configured an API key before.

33. Puzzle accuracy and valid solutions

Puzzles and curated practice should normally have one clearly correct move, or at most two when both are genuinely valid. Two accepted moves should only exist when both are actually correct for the exact position and objective.

Because this is a practice/learning system, every puzzle must have an accurate position, a clearly defined objective and correctly validated solutions. The app must not accept an incorrect move as correct or reject a genuinely correct alternative.

If a puzzle intentionally has more than one valid solution, that should be because the chess position genuinely supports those alternatives — not because the validation is loose. Every accepted move should be verified against

the intended solution logic.

34. AI pop-up on Chess.com — background behaviour

The AI pop-up on the Chess.com home screen needs to be tested to determine whether it actually requires the user's device, app or browser to remain open in the background. Do not assume that a delayed voice response proves that background execution is required.

The delay could instead relate to browser restrictions, audio permissions or another implementation detail. The actual behaviour should be verified and, if background operation is required, the user should be told clearly.

35. AI pop-up — opening and voice delay

The AI pop-up did not open instantly. The user had to click it, and the voice also started slightly late. Investigate whether the click is required by browser audio-permission rules or whether the interaction/loading delay is unnecessary.

The intended experience should feel responsive. If a first user interaction is technically required to unlock audio, handle it gracefully rather than making the feature appear broken or unresponsive.

36. Zara's speaking speed

Zara speaks too quickly. Slow her down so her instructional guidance can be understood comfortably. This should be treated as a specific voice-setting requirement for Zara, alongside the earlier speed feedback for Arjun and Maya.

37. API-key nudges must reflect actual key state

The small message "Want real replies? Add a free AI key." should disappear as soon as a valid AI key has actually been configured. This rule should apply to every API-key nudge throughout the app, not only that one message.

If the user later removes the AI key, the relevant setup prompts should return automatically. In other words: key present means no unnecessary API-key nudge; key absent means the appropriate setup/nudge can appear.

This should be driven by the actual current configuration state, not by a stale UI flag or a setting that remains marked as configured after the key has been removed.

38. Guided "Open in Nightwise / Open in Chess.com / Practise here" shortcuts

Shortcuts such as "Open in Nightwise", "Open in Chess.com" and "Practise here" should not simply present a menu of possible destinations. The course/coach should recommend the best place to practise the particular problem.

The recommendation should be spoken aloud, for example: "I'd recommend opening this in Chess.com" or "I'd recommend practising this here first." The exact recommendation should depend on the learning objective.

The other options can still be available, but the app should clearly identify the recommended option and briefly explain why it is the better place to practise that specific skill.

39. Post-puzzle Arjun feedback

Puzzle navigation was good and the coach conversation itself was good. However, after finishing one puzzle, Arjun did not say anything. There should be an appropriate post-puzzle coaching response.

The feedback should be tied to the actual puzzle and objective: explain what the successful move achieved, why it worked and what lesson the user should carry into a real game. For example, in a pawn-promotion puzzle, Arjun could explain the idea of safely advancing the pawn and promoting it to a new piece.

Do not use a generic post-puzzle message that does not correspond to the actual position.

40. Context-aware greetings and session introductions

Greetings and introductions have repeatedly been a problem. Welcoming a user is an art: the coach needs to know when a welcome is appropriate rather than simply greeting the user every time the app or tab is opened.

For example, if the user has already opened a particular tab today, closed it and returns to it again on the same day, the app should not repeatedly give the full welcome. The user is already in that session context.

However, the first time the user opens that tab on a new day, it becomes a new daily session. “New” does not mean that all content or progress is reset; it means that the first meaningful visit to that tab on that day can receive a fresh welcome and contextual introduction.

The first visit to a tab/section on a given day should therefore greet the user and bring them directly into the current situation — where the tab sits, what the current state is and what they are about to do. Reopening the same tab later that day should preserve continuity without repeating the full greeting.

The important trigger is first meaningful visit to that tab/section per day, not every app launch and not every time the user navigates back to the same tab.

41. Latest-game review — actual opponent move vs Stockfish move

In the latest game review, the system presented the Stockfish best move as though it were the move the opponent actually played. This is inaccurate information and is a serious review-quality problem.

The review must strictly distinguish the real game record from engine analysis. If the opponent actually played one move but Stockfish recommends another, the UI must never describe the Stockfish move as the opponent’s move. The actual move should always be recoverable from the game record.

This is especially important because the user is trying to learn from what really happened in their game. Replacing an opponent’s real move with an engine move changes the story of the game and can teach the wrong lesson.

42. Better structure for move analysis

For important positions, the review can become much more useful by separating what actually happened from what could have happened.

For the user’s side, show what the user actually played and classify it appropriately — for example Best, Good, Average, Inaccuracy, Mistake or Blunder. If a stronger move was available, show the better or best move that could have been played.

For the opponent’s side, show the opponent’s actual reply separately. If their reply was inaccurate, show the better or best response they could have found.

The analysis can then explain what the opponent could have taken advantage of and how the game might have changed if the stronger continuation had been played.

If the user actually played the best move available, do not manufacture a “better move”. Explicitly say that the move was already best where appropriate. Likewise, if the opponent’s real move was the best response, make that clear.

The core model should be: reconstruct the real game first, then compare the actual moves with Stockfish’s optimal lines. The user must always be able to distinguish “what actually happened” from “what could have happened”.

Cross-cutting product principles

Accuracy over appearance — Chess analysis, puzzle validation, progress numbers and voice/API status must reflect reality rather than merely looking plausible.

Preserve uncertainty where appropriate — A confusing result is not automatically proof of one particular bug. Investigate the underlying configuration, objective and data before deciding what caused it.

No incomplete or incorrectly configured puzzles — Every puzzle needs a valid position, clear objective and correct move validation. Ambiguity should come only from genuinely valid chess alternatives.

No fake progress — Unique puzzles, first-attempt correctness, attempts and completion state must be tracked independently and consistently.

Context-aware coaching — Coaches should speak because the user has reached a meaningful point in the session, not simply because a screen opened.

Beginner-first guidance — Professional-level information can remain available, but unfamiliar features need clear onboarding and explanations.

Less clutter, not less capability — Keep information available while progressively revealing it so the user is not overwhelmed.

Truthful system state — Voice, API-key, listening, completion and progress indicators must be based on the actual underlying state.

Daily continuity — A new day can create a new session greeting without resetting the user's accumulated progress.

Actionable recommendations — When there are multiple places to practise, recommend the best destination for the specific learning objective rather than leaving the user to choose blindly.

> This document is intended as the working debriefing for the current app review. It deliberately retains implementation questions and verification points instead of turning every observation into a definitive bug claim.